

Explanation of the C++ OpenMP Code (Min, Max, Sum, Average)

Purpose of the code: This C++ program creates a list of 10,000 random numbers and then computes four statistics on them: the minimum value, the maximum value, the total sum, and the average. It does this using **OpenMP** directives to run the loop in parallel on multiple CPU cores. Parallel loops can greatly speed up processing large arrays by using multiple threads.

C++ Syntax and Core Functions

- **#include lines:** The code begins by including standard headers. For example, `<iostream>` allows using `std::cout` to print output, `<vector>` provides the `std::vector` class for dynamic arrays, and `<cstdlib>` (or `<stdlib.h>`) provides functions like `rand()`. OpenMP code also includes `<omp.h>` (though `<omp.h>` is only needed if you call OpenMP functions directly).
- **Vectors:** A `std::vector<int>` is a dynamic array of integers. Unlike fixed-size C-style arrays, a vector can grow or shrink; the syntax `std::vector<int> arr(10000);` creates a vector named `arr` with 10,000 elements. Internally it stores elements in contiguous memory and resizes automatically ¹.
- **Random numbers:** The function `rand()` generates a pseudo-random integer in the range [0, RAND_MAX). For example, `arr[i] = rand() % 100 + 1;` would fill each element with a random number between 1 and 100. The `srand(time(NULL))` call (often included before using `rand()`) seeds the generator based on the current time so you get different random values each run ².
- **std::min and std::max:** These are functions (in `<algorithm>`) that return the smaller or larger of two values. For example, `min_val = std::min(min_val, arr[i]);` would update `min_val` if `arr[i]` is smaller. (Many codes simply use `if (arr[i] < min_val) min_val = arr[i];` for the same effect.)
- **std::cout:** Used to print text and values to the console. For example, `std::cout << "Min: " << min_val << ...` outputs the results.

Initializing the Array with Random Data

The code typically includes a loop to **fill the vector** with random numbers. For example:

```
srand(time(NULL)); // seed the random number generator
for (int i = 0; i < 10000; i++) {
    arr[i] = rand() % 100 + 1; // random int from 1 to 100
}
```

- Here, `srand(time(NULL))` seeds the generator (using `<ctime>`). Then `rand()` produces pseudo-random values in `0..RAND_MAX`, and the `% 100 + 1` forces them into a 1–100 range ². - The loop `for(int i = 0; i < 10000; i++) { ... }` runs through all indices of the vector, assigning each element `arr[i]`.

Parallel Loop with OpenMP

The key part of the code is the **OpenMP parallel loop**:

```
int min_val = arr[0], max_val = arr[0];
long long sum = 0;

#pragma omp parallel for reduction(min: min_val) reduction(max: max_val) reduction(+: sum)
for(int i = 0; i < 10000; i++) {
    if (arr[i] < min_val) min_val = arr[i];
    if (arr[i] > max_val) max_val = arr[i];
    sum += arr[i];
}
```

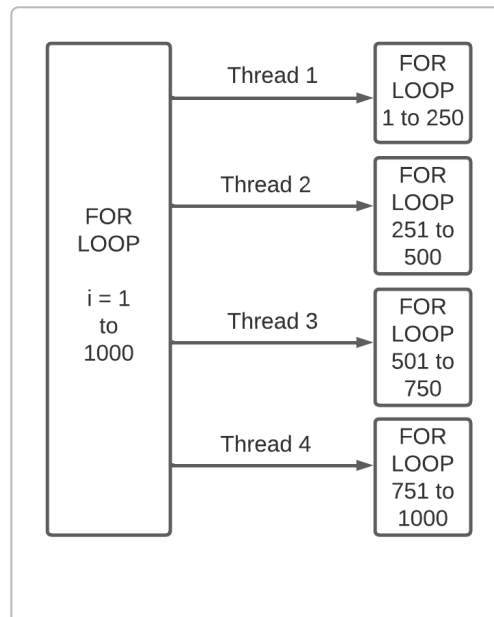


Figure: An OpenMP `parallel for` splits a loop among threads (e.g. Thread 1 handles indices 1–250, Thread 2 handles 251–500, etc.) ³.

- `#pragma omp parallel for`: This directive tells the compiler to run the following loop in parallel. OpenMP will create multiple threads (typically one per CPU core) and split the loop iterations among them ³. For example, if you have 4 threads and 10,000 iterations, each thread might process about 2,500 iterations (as shown in the diagram above). This parallel execution can reduce the total runtime because all threads work at the same time on different parts of the loop.

- **Thread-private copies:** Because multiple threads run the loop body at once, we must avoid them writing to the same variable without coordination. Here, OpenMP's **reduction** clauses handle that. Each thread gets a private copy of the variables `min_val`, `max_val`, and `sum` for its own calculations ⁴. By default, OpenMP initializes each thread's private copy to the appropriate identity value (0 for sum, and a very large or small value for min/max). During the loop, each thread updates **its own** `min_val`, `max_val`, and `sum` without interfering with other threads.
- **Combining results:** After the parallel loop finishes, OpenMP combines the thread-private results into the final values. For example, with `reduction(min: min_val)`, it finds the smallest `min_val` among all threads; with `reduction(max: max_val)`, it finds the largest; and with `reduction(+: sum)`, it sums up all the thread-local sums. This gives the correct global minimum, maximum, and total sum. (OpenMP 3.1 and later defines built-in reductions for arithmetic operations like `+`, and also for `min` and `max` ⁵.)
- **Loop body:** Inside the loop, the code typically does:
 - `if (arr[i] < min_val) min_val = arr[i];` to update the local minimum.
 - `if (arr[i] > max_val) max_val = arr[i];` to update the local maximum.
 - `sum += arr[i];` to add the element to the local sum.
 Each thread applies these to its assigned slice of the array. Thanks to the reduction clauses, these operations are safe in parallel: there is no race condition because each thread has its own variables until the final combination ⁴.

Calculating the Average

After the parallel loop, we have `sum`, the total of all elements, and we know the array size (10000). The average is computed by dividing the sum by the number of elements, e.g.:

```
double avg = static_cast<double>(sum) / arr.size();
```

- We cast `sum` to `double` to perform floating-point division (otherwise, integer division would truncate).
- Then `avg` holds the arithmetic mean of the array values. This is straightforward: "average = total / count."

Why Parallel Code Can Be Faster

- **Multi-core processing:** Modern computers have multiple CPU cores (a Symmetric Multiprocessing system ⁶). OpenMP takes advantage of this by running the loop on several cores at once. Instead of one core doing all 10,000 iterations in sequence, each core/thread does its share (as the figure shows) ³.
- **Work-sharing:** In a non-parallel (serial) loop, each iteration waits for the previous one. OpenMP's `parallel for` splits the iterations so multiple iterations run truly concurrently, reducing overall time ³.
- **Reduction of bottlenecks:** Because each thread accumulates its own sum/min/max, we avoid locking or synchronization inside the loop, which could slow things down. OpenMP's reduction clause handles the combination step efficiently after the loop ends.

- **Performance note:** In practice, the speedup depends on the number of available cores and the work per iteration. For a large array of 10,000 elements, using parallel threads generally gives a big improvement over a single-threaded loop, as long as the overhead of managing threads is small relative to the loop work.

Compiling and Running with OpenMP

To use OpenMP directives, you must enable OpenMP support when compiling. For example, with GCC or Clang, you would do:

```
g++ -fopenmp mycode.cpp -o myprogram
```

The flag `-fopenmp` tells the compiler to recognize `#pragma omp` and link in the OpenMP library ⁷. When you run the program (e.g. `./myprogram`), by default OpenMP will use as many threads as there are CPU cores ⁸. You can control the number of threads with the `OMP_NUM_THREADS` environment variable if desired.

In summary, this program uses simple C++ constructs (`std::vector`, `rand()`, `std::cout`, etc.) for basic tasks, and then uses `#pragma omp parallel for` with reduction clauses to safely run a loop across multiple threads. The result is a faster computation of the min, max, sum, and average on a large array, leveraging multi-core parallelism ⁶ ³.

¹ **Vector in C++ STL | GeeksforGeeks**

<https://www.geeksforgeeks.org/vector-in-cpp-stl/>

² **rand() and srand() in C++ | GeeksforGeeks**

<https://www.geeksforgeeks.org/rand-and-srand-in-cpp/>

³ **C - Parallel for loop in OpenMP | GeeksforGeeks**

<https://www.geeksforgeeks.org/c-parallel-for-loop-in-openmp/>

⁴ **How to: Convert an OpenMP Loop that Uses a Reduction Variable to Use the Concurrency Runtime | Microsoft Learn**

<https://learn.microsoft.com/en-us/cpp/parallel/concrt/convert-an-openmp-loop-that-uses-a-reduction-variable?view=msvc-170>

⁵ **theartofhpc.com**

<https://theartofhpc.com/pcse/omp-reduction.html>

⁶ ⁷ ⁸ **Parallel Programming: Multithreading (OpenMP)**

https://wvuhpc.github.io/2018-Lesson_4/03-openmp/index.html