

Letter Recognition with Neural Networks: A Step-by-Step Guide

In this guide, we walk through Python code for recognizing letters with a neural network. We explain each part of the code in simple terms and include comments and examples.

Imports and Libraries

We start by importing libraries. Each library has a special role:

- **pandas** (`import pandas as pd`): A library for handling data tables. It makes it easy to load, explore, and manipulate datasets ¹.
- **NumPy** (`import numpy as np`): Provides fast array and math operations. It's the base for numerical computing in Python ².
- **scikit-learn** (from `sklearn`): A collection of tools for machine learning. We use it for preprocessing and splitting data (e.g. `StandardScaler`, `LabelEncoder`, `train_test_split`, `confusion_matrix`, `classification_report`) ³ ⁴.
- **TensorFlow/Keras** (`import tensorflow as tf` and `from tensorflow.keras...`): A deep learning library. We use Keras (which is part of TensorFlow) to build and train our neural network ⁵ ⁶. Keras provides easy classes like `Sequential` for models and `Dense` for layers.
- **seaborn** (`import seaborn as sns`): A visualization library built on Matplotlib. It simplifies creating attractive statistical plots, such as a heatmap of a confusion matrix ⁷.

Each import helps organize tasks: Pandas for data loading, NumPy for math, sklearn for data prep and evaluation, TensorFlow/Keras for building the network, and Seaborn for nice plots.

```
import pandas as pd          # data handling
import numpy as np           # numerical computing with arrays
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import seaborn as sns        # for visualization (e.g., heatmaps)
import matplotlib.pyplot as plt
```

Data Loading

We load the letter recognition dataset into a pandas DataFrame. In our code, we use a URL from the UCI Machine Learning Repository, which contains the data in CSV format. We also provide column names because the file has no header row. For example:

```
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/letter-recognition/letter-recogn
col_names = [
    'letter', 'x_box', 'y_box', 'width', 'height', 'onpix',
    'x_bar', 'y_bar', 'x2bar', 'y2bar', 'xybar', 'x2ybar',
    'xy2bar', 'x_ege', 'xegvy', 'y_ege', 'yegvx'
]
data = pd.read_csv(url, header=None, names=col_names)
```

Here, `url` points to the raw data file. We set `header=None` since the file has no header line, and `names=col_names` gives our own column names. The first column `letter` is the actual letter label (A-Z), and the other 16 columns are numeric features describing each letter's shape ⁸ ⁹. For example, `width` and `height` are dimensions of the letter's bounding box, `onpix` is the count of pixels, and so on ¹⁰ ¹¹. In total the dataset has 20,000 samples and 17 columns (16 features + 1 label) ¹² ⁸.

After loading, you can inspect the data briefly (e.g., `data.head()`) to confirm it looks right.

Data Preparation

Next we prepare the data for the neural network. This involves several steps:

- **Separate features and labels:** We split the DataFrame into **X** (inputs) and **y** (targets). The features **X** are all columns except `'letter'`, and the labels **y** are the `'letter'` column (the actual letter for each sample).

```
X = data.drop('letter', axis=1) # all columns except the label
y = data['letter']             # the letter (A-Z) to predict
```

This is just like separating the questions (features) from the answers (labels).

- **Standardize features:** Machine learning algorithms often work better if features are on the same scale. We use `StandardScaler` to make each feature have mean 0 and variance 1 ⁴. In other words, we shift and scale each column so its values are roughly between -1 and 1. This is like normalizing everyone's scores so they're comparable, which helps the neural network learn efficiently.

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Here, `.fit_transform(X)` computes the mean and standard deviation on the training data and scales it. After this, each column of `X_scaled` has average 0 and standard deviation 1. This is a common preprocessing step ⁴.

- **Encode labels:** The letters A-Z are categorical. We convert them to numbers using `LabelEncoder` ¹³. This will map each letter to a number 0–25. For example, 'A'→0, 'B'→1, ..., 'Z'→25.

```
encoder = LabelEncoder()
y_num = encoder.fit_transform(y)
```

Now `y_num` is a numeric array of the letter IDs (0–25) ¹³.

- **One-hot encode labels:** For a neural network output, we usually want a *one-hot* representation: a 26-length binary vector where only the correct letter's position is 1. Keras provides `to_categorical` for this.

```
y_onehot = tf.keras.utils.to_categorical(y_num)
```

This turns each numeric label like 3 into a vector like `[0,0,0,1,0,...,0]`. One-hot encoding creates separate “dummy” columns for each class ¹⁴. In our case, `y_onehot` will be a 20,000×26 array of 0s and 1s, where each row has a 1 in the column for that letter.

In summary, data preparation separates inputs and outputs, scales inputs to a common range, and turns outputs into numeric form (then one-hot vectors) for the network.

Train-Test Split

Before training the model, we split the data into **training** and **testing** sets. This ensures we can evaluate the model on “new” data it hasn't seen during training. In code:

```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_onehot, test_size=0.2, random_state=42
)
```

This shuffles the data and keeps 20% for testing (`test_size=0.2`) and 80% for training ¹⁵. The `random_state` makes the split reproducible. This process is called *train-test split*. It is a common **validation technique**: the model is trained on the training set and then its performance is measured on the test set to simulate how it would do on new, unseen data ¹⁵ ¹⁶. For example, one source explains that the

training set is used to fit the model while the testing set (which the model hasn't seen) is used to check its performance ¹⁶ .

Model Building

Now we build the neural network. We use Keras's **Sequential** model, which is just a linear stack of layers ⁶ . Each layer is a *Dense* layer (fully connected to the previous layer). In code:

```
model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    Dense(26, activation='softmax')
])
```

- The first **Dense** layer has 64 neurons and uses the **ReLU** activation function. ReLU (Rectified Linear Unit) sets negative inputs to zero and keeps positive inputs as is ¹⁷ . This non-linearity lets the network learn complex patterns; it helps the model learn faster by avoiding the vanishing gradient problem ¹⁷ . You can think of ReLU like a gate that only lets positive signals through.
- The second (and final) **Dense** layer has 26 neurons (one for each letter A-Z) and uses **Softmax** activation. Softmax takes the raw outputs (called logits) and converts them into probabilities that sum to 1 ¹⁸ . In effect, the network outputs 26 values that can be interpreted as probabilities for each letter. The highest probability is the network's guess. Softmax is the standard choice for multi-class classification because it turns outputs into a probability distribution ¹⁸ .

Each **Dense** layer applies a weight matrix to its inputs and adds a bias, then applies the activation function. The `input_shape=(X_train.shape[1],)` tells Keras the number of input features (16 in this case).

By stacking these layers, the model will learn to map the 16 input features to one of 26 output letters. The ReLU layer learns intermediate features, and the Softmax layer produces the final letter probabilities. This simple architecture (one hidden layer) is often enough for this task, but more layers or neurons could be used for more complexity.

Model Compilation and Training

With the model defined, we **compile** it by choosing a loss function, optimizer, and metrics:

```
model.compile(
    loss='categorical_crossentropy',
    optimizer='adam',
```

```
metrics=['accuracy']  
)
```

- **Loss function:** We use **categorical_crossentropy** for multi-class classification. This loss measures the difference between the true one-hot labels and the predicted probabilities. It is often called *log loss*. In essence, crossentropy will be small when the predicted probability for the correct class is high, and large when the model is very wrong. It “guides” the model to improve its predictions ¹⁹.
- **Optimizer:** We use **Adam**. Adam is an adaptive gradient optimizer that adjusts learning rates for each weight individually. It is known for fast convergence and works well in many deep learning tasks ²⁰. In simple terms, Adam changes the model's weights in a smart way to minimize the loss efficiently.
- **Metrics:** We include `'accuracy'`. During training and evaluation, Keras will compute accuracy (the fraction of correct predictions) alongside the loss.

Next, we train (`fit`) the model:

```
history = model.fit(  
    X_train, y_train,  
    epochs=20,  
    batch_size=32,  
    validation_split=0.1  
)
```

- **Epochs:** Each epoch means one pass through the entire training set. Here we train for 20 epochs. More epochs usually means the model can learn better, up to a point ²¹. (Each epoch sees all 16,000 training samples once.)
- **Batch size:** The number of samples the model looks at before updating weights. A batch of 32 means the model processes 32 samples, computes the loss, and updates weights, then moves on to the next 32, and so on ²². Smaller batches give noisier but more frequent updates; larger batches are smoother but use more memory.
- **Validation split:** We set `validation_split=0.1`, which means 10% of the *training* data is held out for validation each epoch. The model will report loss and accuracy on this validation set, which helps us see if it's learning well without overfitting ²³.

All these settings are hyperparameters you can tune. In summary, compilation and training tell Keras *how* to learn (loss & optimizer) and *how long/what batch size* to use.

Model Evaluation

After training, we evaluate the model on the test set to measure performance. For example:

```
loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test accuracy: {accuracy:.3f}")
```

- **Loss:** The value of the loss function (crossentropy) on the test data. Lower is better, meaning the predicted probabilities match the true labels more closely.
- **Accuracy:** The fraction of test samples the model got right. It is computed as (number of correct predictions) / (total predictions). In classification, accuracy tells us the proportion of all predictions that were correct ²⁴. For instance, an accuracy of 0.90 means 90% of letters were classified correctly.

Thus, we assess the model by its test loss and accuracy. High accuracy and low loss indicate a good model.

Prediction and Results Analysis

Finally, we use the model to make predictions and analyze the results. We typically convert the model's outputs back into letter labels and inspect detailed metrics:

- **Predictions to labels:** The model's `predict` method will output a probability vector for each test sample. We use `np.argmax` to pick the index of the highest probability for each sample, which gives the predicted class number. Then we can map these numbers back to letters using the `LabelEncoder`:

```
y_pred_probs = model.predict(X_test)
y_pred_classes = np.argmax(y_pred_probs, axis=1)
y_true_classes = np.argmax(y_test, axis=1) # since y_test is one-hot
# Optionally convert back to letters:
pred_letters = encoder.inverse_transform(y_pred_classes)
true_letters = encoder.inverse_transform(y_true_classes)
```

- **Confusion Matrix:** A confusion matrix is a table that compares actual labels to predicted labels. It lets us see *where* the model is making mistakes ²⁵. Each row of the matrix represents the actual class, and each column represents the predicted class. For a simple 2-class example, the matrix shows True Positives, False Positives, False Negatives, and True Negatives. Here's an illustrative example (blue and red colors are just for visual emphasis):

| | | Actual | |
|-----------|----------|----------------|----------------|
| | | Positive | Negative |
| Predicted | Positive | True Positive | False Positive |
| | Negative | False Negative | True Negative |

Example confusion matrix: actual vs. predicted labels (Positive/Negative).

In a multi-class setting like our 26 letters, the confusion matrix would be 26×26. A cell [i,j] in the matrix shows how many samples of actual class i were predicted as class j. We can compute it in code with scikit-learn's `confusion_matrix`:

```
cm = confusion_matrix(y_true_classes, y_pred_classes)
```

We can then visualize it with a heatmap (using seaborn) to spot patterns (dark diagonal means correct predictions, off-diagonals show confusions). A confusion matrix highlights which specific letters are often mistaken for others ²⁵.

- **Classification Report:** Another summary is the *classification report*, which includes precision, recall, and F1-score for each class. In scikit-learn we can do:

```
report = classification_report(
    y_true_classes, y_pred_classes,
    target_names=encoder.classes_
)
print(report)
```

- **Precision** for a class is the fraction of predictions for that class that were correct: $TP / (TP + FP)$. It tells us how “pure” the positive predictions are ²⁶ ²⁷.
- **Recall** is the fraction of actual class examples that were correctly identified: $TP / (TP + FN)$ ²⁶ ²⁷. It shows how many of the true letters we caught.
- **F1-score** is the harmonic mean of precision and recall ²⁸. It balances both aspects into one number (best is 1.0, worst is 0.0).

For example, an F1 of 0.95 for letter 'A' means the model has high precision and recall on 'A'. These metrics help us understand performance per letter, beyond overall accuracy ²⁶ ²⁸ .

In summary, after training we predict on the test set, then analyze results with tools like the confusion matrix and classification report. These tell us not just the overall accuracy, but *how* the model is performing on each letter (precision, recall, etc.), which is crucial for a thorough evaluation.

Each of the above steps ties directly to the code components in the script, providing a clear picture of how the neural network letter recognizer works from start to finish.

Sources: Authoritative documentation and tutorials on pandas, NumPy, scikit-learn, Keras, and data science practices were used to explain these concepts ¹ ² ³ ⁵ ⁶ ⁴ ¹³ ¹⁴ ¹⁵ ¹⁶ ¹⁷ ¹⁸ ¹⁹ ²⁰ ²² ²³ ²⁴ ²⁵ ²⁶ ²⁸ . Each citation supports the descriptions above.

-
- ¹ **pandas - Python Data Analysis Library**
<https://pandas.pydata.org/>
 - ² **NumPy**
<https://numpy.org/>
 - ³ **Learning Model Building in Scikit-learn | GeeksforGeeks**
<https://www.geeksforgeeks.org/learning-model-building-scikit-learn-python-machine-learning-library/>
 - ⁴ **StandardScaler — scikit-learn 1.6.1 documentation**
<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>
 - ⁵ ⁶ **Keras: The high-level API for TensorFlow | TensorFlow Core**
<https://www.tensorflow.org/guide/keras>
 - ⁷ **seaborn: statistical data visualization — seaborn 0.13.2 documentation**
<https://seaborn.pydata.org/>
 - ⁸ ⁹ ¹⁰ **UCI Machine Learning Repository**
¹¹ <https://archive.ics.uci.edu/ml/datasets/letter+recognition>
 - ¹² **Letter Recognition Classification Problem | by Johannes Aniciete | Medium**
<https://medium.com/@mjfaniciete/letter-recognition-classification-problem-ce91a169935c>
 - ¹³ **LabelEncoder — scikit-learn 1.6.1 documentation**
<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.LabelEncoder.html>
 - ¹⁴ **Python Keras | keras.utils.to_categorical() | GeeksforGeeks**
https://www.geeksforgeeks.org/python-keras-keras-utils-to_categorical/
 - ¹⁵ **train_test_split — scikit-learn 1.6.1 documentation**
https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html
 - ¹⁶ **Train Test Split: What it Means and How to Use It | Built In**
<https://builtin.com/data-science/train-test-split>
 - ¹⁷ **A Gentle Introduction to the Rectified Linear Unit (ReLU) - MachineLearningMastery.com**
<https://machinelearningmastery.com/rectified-linear-activation-function-for-deep-learning-neural-networks/>

- 18 **Softmax Activation Function in Neural Networks | GeeksforGeeks**
<https://www.geeksforgeeks.org/the-role-of-softmax-in-neural-networks-detailed-explanation-and-applications/>
- 19 **Categorical Cross-Entropy in Multi-Class Classification | GeeksforGeeks**
<https://www.geeksforgeeks.org/categorical-cross-entropy-in-multi-class-classification/>
- 20 **Complete Guide to the Adam Optimization Algorithm | Built In**
<https://builtin.com/machine-learning/adam-optimization>
- 21 22 **How to choose Batch Size and Number of Epochs When Fitting a Model? | GeeksforGeeks**
<https://www.geeksforgeeks.org/how-to-choose-batch-size-and-number-of-epochs-when-fitting-a-model/>
- 23 **Model training APIs**
https://keras.io/api/models/model_training_apis/
- 24 **Classification: Accuracy, recall, precision, and related metrics | Machine Learning | Google for Developers**
<https://developers.google.com/machine-learning/crash-course/classification/accuracy-precision-recall>
- 25 **Understanding the Confusion Matrix in Machine Learning | GeeksforGeeks**
<https://www.geeksforgeeks.org/confusion-matrix-machine-learning/>
- 26 27 28 **Machine Learning: Precision, Recall, F1-Score | by Mudgalvivek | Medium**
<https://medium.com/@mudgalvivek2911/machine-learning-precision-recall-f1-score-c0a064ea6008>